

**Manage Big Data
in the Cloud
-- Different NoSQL DBs**



MongoDB

❖ NoSql systems

- Most of those in the wide column family and key-value store categories have similar designs
- Document based NoSql has some new data model concept
 - Similar to old time “object-oriented” database model

❖ Applications

- Content management systems
- Product management systems
- ...

MongoDB – Data Model

❖ Semantics

- MongoDB = a set of databases
 - Database = a set of collections
 - Collection = a set of documents
 - With similar data fields, but still highly flexible
 - Documents
 - A set of attributes and their values
 - Use BSON format (Binary JSON)
 - Support nesting: Documents can include documents
 - Document example (nested):
 - { “name” : “Alice Anderson”, “SSN” : “123456789”,
“member” : “UTD”, “member” : “IEEE”,
“address” : { “street” : “800 W. Campbell”,
“city” : “Richardson”, “state” : “TX” } }
Can have any #values for the same field
- Can have nested documents
⇒ Offer better semantics

MongoDB – Data Model

❖ Semantics

➤ Relationship

■ Embedded relationship

- { “name” : “Alice Anderson”, “SSN” : “123456789”, “address” : [{ “street” : “a”, “city” : “b”, “state” : “c” } { “street” : “d”, “city” : “e”, “state” : “f” }] }
- Each entry in [] is a document, [] embeds the entries

■ Reference relationship

- Similar to RDB foreign key
- { “Id” : ObjectId(“12345”), “name” : “Alice Anderson”, “SSN” : “...”, “address” : [ObjectId(“13579”), ObjectId(“24680”)] }
- Objects within the relationship definition are first class objects

MongoDB – Data Model

❖ Query

➤ Can lookup by any keyword defined in the document

- E.g., db.users.find ({ “name” : “Alice”, “age” : { “\$gte”, 18 } })

➤ Can specify what fields to return

- E.g., db.users.find ({ ... }, { “email”:1, “birth-date”:1 })
- 0 means not to return the corresponding field

➤ Can specify the subfields for lookup

- E.g., db.users.find ({ “address” : { “state” : “Texas” } })

users						
name	email	phone	address	age	birth	...

MongoDB – Data Model

❖ Indexing

- Index is created for efficient search
 - Data structure: B-tree
- User can specify the field to be indexed
 - E.g., `db.users.ensureIndex ({"name":1, "age":1})`
- Can index the subfields also
 - E.g., `db.users.ensureIndex ({"address.state":1})`

❖ Indexes

- Cached in memory and stored on disk
- Stored in a shard unit


MongoDB – Consistency

❖ Support multiple consistency models

- But not as many as Cassandra
- Passive replication with primary and backup
- Strong consistency for writes
 - Can have a set of primary nodes (primary set)
 - Update is sent to all nodes in the primary set
 - If a primary fails, a secondary can be elected to be the primary
 - Atomicity is guaranteed only for an individual document, not cross documents
- Reads
 - If strong consistency is desired, then read from the primary
 - If only require eventual consistency, then read from any

MongoDB – Sharding

- ❖ Range based sharding
 - Similar to GBT
 - Each chunk contains a key range
- ❖ Centralized shard distribution
 - Mongo is the query router
 - Config server stores the centralized metadata and configuration settings for the system



Graph Databases

❖ Why graph database

➤ Applications

- The web, the Linked Data, social networks, etc.
- The RDF database
- Many data are highly linked: bio data, chemical data, ...

➤ Why are relational and nosql databases not sufficient?

- They can capture all the relations
- But query processing can be very slow
- E.g., in a social network database
 - Find all Alice's friends: can be done efficiently in conventional DBs
 - Find all people having Alice as a friend: slow query processing

Graph Database Models


❖ Neo4j

<https://neo4j.com/docs/getting-started/current/>
The community version is free

➤ Model: (Nodes, relationships) ← (labels and properties)

➤ Example

- Node with properties: (name:Alice, email:a@utd.edu, ...)
- label for nodes: person A node can have multiple labels
 - Labels can support grouping, etc.
- Relationships can also have properties

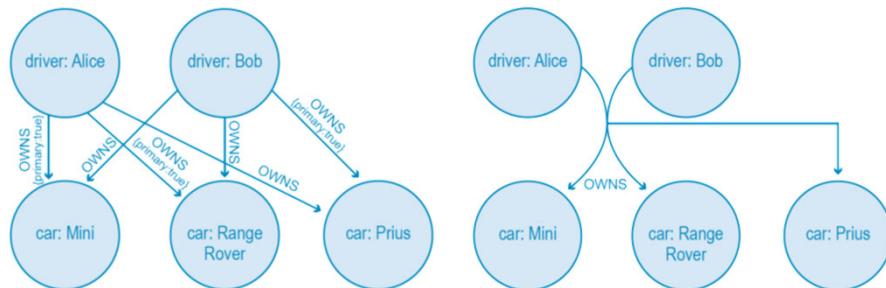


From: Graph Databases: New Opportunities for Connected Data

Graph Database Models

❖ General graph database models

- Simple graph model: (starting node, relation, ending node)
 - E.g., Neo4j
- Hypergraph model: M to N relation
 - A couple owns three cars \Rightarrow 6 relations \Rightarrow 1 hyper-relation

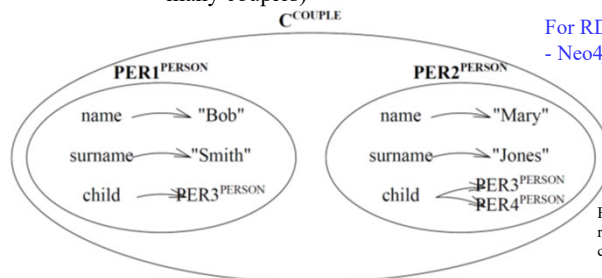


From: Graph Databases: New Opportunities for Connected Data

Graph Database Models

❖ General graph database models

- Nested graph model
 - Each node may also be a graph
 - Example 1
 - Neo4j: use properties for each person (better)
 - » Add couple partner as a property
 - Neo4j: use labels for couples (may not be ideal when there are too many couples)



For RDF model \Rightarrow nesting helps!
- Neo4j does criticize RDF model

From: A nested-graph model for the representation and manipulation of complex objects

Graph Database Models

❖ General graph database models

➤ Nested graph model

For this example, nesting is better!

■ Example 2: Activities in a company

- Inter-division, company wide
 - » E.g., benefit selection to everyone, company wide gatherings
 - » E.g., hiring request from engineering to HR and back
- Inter-group in a division
 - » E.g., in the software engineering group: system group with database group (relation: design meeting, software delivery), database group with QA group (relation: delivery for testing)
- Intra group
 - » E.g., in the software engineering group: group meetings, manager to group member evaluations, member to member discussions for design, code integration, etc.
- Other
 - » Individuals cross group/division activities

If no grouping, how many edges will there be?

Graph Database Queries

❖ Neo4j

- <https://neo4j.com/docs/cypher-manual/current/clauses/match/>
- Graph Databases: New Opportunities for Connected Data

➤ Match and Return clauses

➤ Node finding

- MATCH (movie:Movie) RETURN movie.title
 - Return movie.title field from all nodes with label "Movie"
- MATCH (director {name: 'xxx'})--(movie) RETURN movie.title
 - Return movie.title from all nodes with property "director.name=xxx"
- Node finding can be handled similarly with relational database

➤ Relation finding

- MATCH (:Person {name: 'xxx'})-[r]->(movie) RETURN type(r)
 - Returns the types of the outgoing relations from "person.name=xxx"
 - E.g., acted, directed

Graph Database Queries

❖ Neo4j

➤ Path based match

- `MATCH (a:Person)-[:Knows]->(b)-[:Friend]->(c),
WHERE a.name = 'xxx' RETURN b, c`
 - Returns all b and c where b knows xxx and c is a friend of b
- `MATCH (Movie {title: 'xxx'})-[*0..1]-(x) RETURN x`
 - Returns all movies with title xxx (path length=0) and all the nodes that are 1 hop away with the matching nodes (path length=1)
- `MATCH p = (p1:Person)-[*]-(p2:Person) WHERE p1.name = 'xxx'
AND p2.name = 'yyy' RETURN p`
 - Returns the entire path between nodes p1 and p2 with the specific name property (p1.name=xxx, p2.name=yyy)

Graph Database Queries

❖ General graph database queries

➤ Subgraph query

- Search for a specific pattern in a graph
- Some part of the pattern may be uncertain
- Return the graphs that contains the pattern and has the match
- Path based match in Neo4j has some similarities to this

➤ Supergraph query

- Search in a graph database $\{g1, g2, \dots\}$ and find all graphs that are contained in the input graph given by the query

➤ Similarity queries

- Instead of exact match, consider similarity based matching
- Similar in node names/properties, relationship names/properties
- Similar in the graph pattern, not exactly isomorphic graphs

Graph Database Storage

❖ Graph database storage

- Most existing works still use relational DB tables
 - E.g., RDF databases

➤ Neo4j graph storage

- Use native graph storage(?), but still different from graph itself
- Storages for nodes, relationships, properties, labels, ...
- Nodes
 - In use: if not in use (e.g., deleted), can be reclaimed
 - Pointers to properties, labels, relationships (4 bytes)

In-use	Relations ↓	Properties ↓	Labels ↓	flags
--------	-------------	--------------	----------	-------

- Relationships
 - Relationship type pointer: to relationship type table
 - Starting/Ending node pointers: to node table
 - Doubly linked list (previous/next relation for the node)

Relation type ↓	Start node ↓	End node ↓	pre ↓	next ↓
-----------------	--------------	------------	-------	--------

Graph Database Storage

❖ Graph database storage

➤ Neo4j graph storage

- Properties
 - Each record can have 4 properties
 - » Each is a key-value entry
 - » Each have inline property value or a pointer to property value store (if the value is large)
 - Singly linked list to next properties of the same node (> 4 properties)
- Labels: no discussion, can be similar to property
- Node, relationship, property stores: have fixed sized records
 - Pointers are like the array indexes to the records in the table (ids)
- Match: $n: \{property\} - r: \{property\} \rightarrow m$
 - Property table: find the property id specified for n and r (starting id)
 - Node table: find the node ids of nodes with the property id
 - Relation table: find matching starting node ids and relationship ids

Type ↓	Key	Value (↓)	next
Type ↓	Key	Value (↓)	↓
Type ↓	Key	Value (↓)	
Type ↓	Key	Value (↓)	

Graph Database Indexing Techniques

❖ Neo4j

- Some tables are like indexes

❖ Other, general

- Tables for nodes
 - With detailed properties and labels
 - With incoming and outgoing edges
 - Each edge can have corresponding node pointers (other nodes)
- Indexing: Keyword based inverted indexing
 - A table with all keywords in the database
 - Points to nodes (in any field)
- Pattern matching
 - Match all keywords for nodes and find matching nodes
 - From these nodes, find matching relationships give in the pattern

Distributed Graph Databases

❖ Neo4j

<https://neo4j.com/docs/operations-manual/current/clustering/introduction/>

- Only available in the enterprise version
- Mainly consider replication for availability
 - A cluster may host multiple graphs, which can be hosted by different nodes, but a single graph is not partitioned
 - Still investigating partitioning techniques for a single graph
 - Primary and secondary hosts
- Consistency
 - Consider causal consistency
 - A client should see the latest updates it knows about
 - Client: issues an update → gets back a **bookmark** → issues subsequent r/w requests with the bookmark
 - Server: processes the r/w request only if the bookmarked update has been completed

Distributed Graph Databases

❖ Pregel

- MapReduce on graph databases, by Google
- Partitioning
 - Graph DB is a collection of vertices, each with a unique ID
 - Each vertex is hashed to a node by $\text{hash}(\text{ID}) \% N$
 - N is the number of storage nodes
- Computation
 - Implement supports for functions that are similar to MapReduce
 - Pass messages (shuffling)
 - Aggregation (reduce)
 - ...

Distributed Graph Databases

❖ Sedge

- Mainly focuses on graph database partitioning
- Computation steps are based on Pregel
- Partitioning
 - Principle: minimize communication cost during query processing
 - Partitioned by graph minimal cut
 - Repartitioning
 - When there are many cross-partition accesses \Rightarrow Repartition
 - Partitioning is evolving continuously
- Two level partition management
 - Primary partitioning: partitions are disjoint
 - One vertex can only belong to one partition
 - Secondary partitioning: balance workloads

Distributed Graph Databases

❖ Sedge

➤ On demand partitioning

- Manage partitions to achieve the secondary partitioning principle
- Partition replication
 - When a partition becomes a hotspot, replicate it
 - » Should be referring to query processing without updates (write to all, so no benefit to replicate)
- Dynamic partitioning
 - Decide new optimal partitioning based on historical queries
 - » Without considering original partitioning
 - Apply the new partitioning result
 - » ??? Too much data migration

Distributed Graph Databases

❖ Sedge

➤ Dynamic partitioning

- Need to track query processing traces (vertices and edges visited)
⇒ Too much information to track
- Group vertices in k hops (k=1 has the best effect)
 - Partition the graph into disjoint groups, radius of each group is k
 - Choose one vertex, put all its neighbors within k hops in a group
 - Give the group a color (or just an ID)
- Query history
 - Track the colors that each query traversed
 - All the colors of a query form an envelop
 - » E.g., query Q1, Q2 has colors {c1, c3, c5}, {c1, c2, c8}
 - Recording visited groups instead of vertices ⇒ space/time efficient

Distributed Graph Databases

❖ Sedge

➤ Dynamic partitioning

- Cluster the envelopes (queries) based on their similarity
 - Jaccard similarity = $|Q1 \cap Q2| / |Q1 \cup Q2|$
 - E.g., similarity of {c1, c2, c3, c5} and {c1, c3, c5, c7} is 3/5
 - \Rightarrow Each cluster includes a set of colors, which is the union of all the colors of all the queries in the cluster
- Rank the clusters by $\rho = (\text{\#queries}) / (\text{\#colors})$
 - Access frequency for each color
 - Sort the clusters by ρ
- Greedy algorithm for partitioning Pack similar clusters \Rightarrow More groups, less vertices
 \Rightarrow Because similar groups has more overlaps
 - Each partition is of size S (predetermined)
 - Pack as many clusters (largest ρ first) to the current partition
 - » I.e., add the colors (groups of vertices) in the cluster to the partition
 - Switch to the next partition when the current partition is filled

References

➤ MongoDB

- Kristina Chodorow, Michael Dirolf. “MongoDB: The Definitive Guide”, O’Reilly Media, Sep. 2010

➤ Graph DBs

- Graph Databases: New Opportunities for Connected Data (Neo4j)
- Pregel: A system for large-scale graph processing (Pregel)
- Towards effective partition management for large graphs (Sedge)
- An introduction to graph data management (models & queries)
- Graph indexing and querying: A review (indexing techniques)