

# Cloud Storage using Erasure Code

## Distributed Storage

### ❖ Striping/Sharding


- How to store objects in distributed storage
  - A lot of data objects (or files) in many storage nodes
  - A large data object in many storage nodes  $\Rightarrow$  Striping
    - We have discussed how to partition, placing, maintaining layout
- What are the problems with striping?
  - One node failure  $\Rightarrow$  data loss
    - One node failure:  $f \Rightarrow$  At least one of  $N$  node failure:  $1 - (1-f)^N$
- Solutions
  - Striping + replication  $\Rightarrow$  High space consumption
  - $(N,K)$  erasure code
    - Divide data into  $N$  shares, reconstruct data from  $K$  shares

## Erasure Code

### ❖ (N,K) erasure code


#### ➤ Code with original data

- RAID is also a type of erasure code with original data



#### ➤ Code without original data

- May be for security, e.g. secret sharing
- E.g.,  $N = 10, K = 4$ , take any 4, original data can be reconstructed



## Erasure Code

### ❖ Different focuses for communication and storage

- Most erasure codes are originally designed for reliable data transmission
- In communication
  - Error checking and recovery are both important
- In storage
  - Consider storage node failures, but failure detection is not the main concern, only need to consider recovery from failures
    - Each piece of data is stored at a different storage  $\Rightarrow$  Failed node is clearly known
    - Many disks have their own integrity coding schemes to detect errors

## Erasure Code, Replication, Striping


### ❖ What are to be compared

- Redundancy level
  - How much more redundant space is needed to achieve the same level of fault tolerance
- Parallel transmission
  - When accessing a certain data by one client, different parts of the data can be transmitted from different nodes
  - Reduce the communication latency
- Parallel accesses
  - When accessing the same data by different clients, concurrent client accesses can be parallelized
  - Reduce the potential waiting time

## Erasure Code, Replication, Striping

### ❖ Redundancy level


- Replication
  - Consider 3 replicas: with 2 failures, there is still one good replica
  - Storage space is tripled, redundancy is 200%
- (N, K) erasure code
  - 4 data nodes, 2 redundancy nodes
  - Can also tolerate 2 failures
  - Only 50% redundancy
  - Any disadvantage?
    - Recovery time will be higher
- Striping + replication
  - Same redundancy as replication
    - (striping alone has no redundancy, no fault tolerance)



## Erasure Code, Replication, Striping


### ❖ Parallel accesses and communication cost

#### ➤ Striping + replication



- Cannot access arbitrarily
- Has to get one block from each specific node group

#### ➤ (N,K) erasure code



- Assume MDS: Maximum distance separable code
- Data can be reconstructed from any K (out of N) shares
- Increase N to increase the level of parallel accesses
- Access the nearest K nodes to reduce the communication cost
- The data encoding and reconstruction cost is a concern

## RAID

### ❖ RAID 4

0	1	2	3	P0
4	5	6	7	P4
8	9	10	11	P8
12	13	14	15	P12
Disk 0	Disk 1	Disk 2	Disk 3	Disk 4

0	0	0	0	0
1	0	0	0	1
0	0	1	1	0
1	1	0	0	0

Read: do not access parity  
Write: need to read/write parity  
⇒ Unbalanced disk load

### ❖ RAID 5

0	1	2	3	P0
4	5	6	P4	7
8	9	P8	10	11
12	P12	13	14	15
P16	16	17	18	19
Disk 0	Disk 1	Disk 2	Disk 3	Disk 4

Distribute parity disk  
⇒ Balanced disk load for both read and write operations

## RAID

### ❖ RAID 6

- 2 parity disks
- How to compute the second parity? Diagonal
  - Consider K data disks and K data size ( $w = K$ )

0	1	2	3	P0	P0
4	5	6	7	P4	P4
8	9	10	11	P8	P8
12	13	14	15	P12	P12
Disk 0	Disk 1	Disk 2	Disk 3	Disk 4	Disk 5
0	0	0	0	0	0
1	0	0	0	1	0
0	0	1	1	0	0
1	1	0	0	0	1

- But recovery may be a problem!

## RAID

### ❖ Recovery problem in simple 2-parity solution

- Consider two disk failures

0	?	0	?	0	0
0	?	0	?	0	0
0	?	0	?	0	0
0	?	0	?	0	0

- Multiple solutions

0	0	0	0	0	0	0	1	0	1	0	0
0	0	0	0	0	0	0	1	0	1	0	0
0	0	0	0	0	0	0	1	0	1	0	0
0	0	0	0	0	0	0	1	0	1	0	0

## EvenOdd Coding

### ❖ EvenOdd coding

#### ➤ Consider K disks, K-1 block size

- K is a prime

0	1	1	0	0	0	1
1	1	0	0	0	1	0
0	1	1	1	0	0	0
0	0	0	1	1	1	0
0	0	0	0	0	0	0

Pseudo row, every bit should be 0  
 (5,7) is diagonal parity for gray  
 Need to make this bit (5,7) = 0  
 ⇒ Let this be the adjust bit A  
 ⇒ For this example, A = 1  
 ⇒ Use A to adjust all the diagonal parity bits

#### ➤ How to compute A from parities?

- XOR (horizontal parities)  
 = XOR (all data bits,  $K * (K-1)$  bits)  
 = XOR (diagonal parities) + A
- In this example:  $\text{XOR}(0110) = \text{XOR}(1000) + A \Rightarrow A = 1$

## EvenOdd Coding

### ❖ Recovery

#### ➤ How to recover from 1 disk failure?

- Data disk failure: recover from one of the parity disks
- Parity disk failure: recompute the lost parity

#### ➤ How to recover from 2 disk failures?

- 1 data disk + diagonal parity disk
  - First recover the data disk from horizontal parity disk
  - Then recompute the diagonal parity disk
- 1 data disk + horizontal parity disk
- 2 data disks

#### ➤ Why do we need to adjust the diagonal parity???

## EvenOdd Coding

### ❖ Recovery from 2 data disk failures

[illegible]

- For any two data disk failures, there are always two diagonals that only lost one bit  $\Rightarrow$  derive those lost bits first
  - Disks 1 and 2 fail: orange bit  $M[1,1]$  can be computed first
  - Disks 3 and 5 fail: blue bit  $M[2,3]$  can be computed first
- After the first recovery, one row can be recovered fully, then zig-zag to recover the rest

## EvenOdd Coding

- ❖ Recovery from 1 data disk+ horizontal parity disk failures

0	1	1	0	0	0	1	1
1	1	0	0	1	0	1	0
0	1	1	1	0	1	1	0
0	0	0	1	1	1	0	0
0	0	0	0	0	0	0	0

- With any 1 data disk failure, there is always one diagonal without data loss  $\Rightarrow$  Derive the adjust bit from that diagonal

## Even-Odd Coding

### ❖ Why N has to be prime?

■ (N: # data disks, e.g., 4, 6)

#### ➤ By counter example

	1	2	3	4		
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
0	0	0	0	0	0	0

	1	2	3	4	5	6		
1	0	0	0	0	0	0	0	0
2	0	1	0	1	0	0	0	1
3	0	0	0	0	0	0	0	0
4	0	1	0	1	0	0	0	0
5	0	0	0	0	0	0	0	1
0	0	0	0	0	0	0	0	0

When #data-disk is even, (N = 4/6):

Choose disks 1,3 as failed (gray, green)

⇒ Gray on disk 3 is at row  $(N-(3-1)) = 2/4$

⇒ Start from row 2 on disk 1, then disk 3,  
next row is always "2" rows after  
i.e., row  $(s*2) \% N$ ,  $s = 1..N-1$

⇒ Always get to gray before finish ⇒ Stuck  
(because gray lost only one bit,  
after recovering the lost bit,  
cannot proceed on diagonal any more)

When #data-disk is multiple of 3 (e.g., 9/15)

Choose disks 1,4 as failed (gray, green)

⇒ Gray on disk 4 is at  $(N+1-4) = 6/12$

⇒ Start from position 3, jump 3

⇒ Always get to gray before finish ⇒ Stuck

## RAID-DP

### ❖ RAID double parity

#### ➤ Used in RAID-6

#### ➤ Consider K disk and K data bits from each disk

■ K+1 is a prime

0	1	1	0	0	1
1	1	0	0	0	0
0	1	1	1	1	1
0	0	0	1	1	1

#### ➤ Horizontal parity

#### ➤ Diagonal parity

■ The horizontal parity is also used to compute the diagonal parity

■ Gray diagonal has no diagonal parity



## RAID-DP

	1	2	3	4	5	
1	0	1	1	0	0	1
2	1	1	0	0	0	0
3	0	1	1	1	1	1
4	0	0	0	1	1	1
0	0	0	0	0	0	0

1 disk failure or  
1 data disk + D-parity disk  
⇒ no problem

### ❖ Any two disk failures, not the diagonal parity

- Horizontal parity disk is the same as data disk
  - They are horizontal parities for each other
- If disk 1 and another disk failed (e.g., 3)
  - Always can find a non-gray diagonal with only 1 lost bit
    - The other diagonal is gray, which also lost only one bit
  - Blue: (2,1) → (2,3) → (4,1) → (4,3) → (1,1) → (1,3) → (3,1) → (3,3) → done

## RAID-DP

	1	2	3	4	5	
1	0	1	1	0	0	1
2	1	1	0	0	0	0
3	0	1	1	1	1	1
4	0	0	0	1	1	1
0	0	0	0	0	0	0


### ❖ Any two disk failures, not the diagonal parity

- If failed does not include disk 1 (e.g., 2 and 5)
  - Always can find 2 non-gray diagonals with only 1 lost bit
  - Orange: (3,2) → (3,5) → (1,2) → (1,5) → stuck
  - Green: (2,5) → (2,2) → (4,5) → (4,2) → stuck, but done

## Reed Solomon Code

### ❖ Generation matrix

- Identify matrix  $I_n$ 
  - With dimension  $n \times n$
- Additional rows: Redundancy matrix
  - Vandermonde and Cauchy constructions Assure invertibility
- In encoding, no need to compute first  $n$  rows
  - They are the original data (because of the identify matrix)



## Reed Solomon Code

### ❖ Vandermonde matrix

$$E = \begin{bmatrix} 1 & \alpha_1 & \alpha_1^2 & \dots & \alpha_1^{n-1} \\ 1 & \alpha_2 & \alpha_2^2 & \dots & \alpha_2^{n-1} \\ 1 & \alpha_3 & \alpha_3^2 & \dots & \alpha_3^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \alpha_m & \alpha_m^2 & \dots & \alpha_m^{n-1} \end{bmatrix}$$

### ❖ Cauchy matrix

$$E^T = a_{ij} = \frac{1}{x_i - y_j}; \quad x_i - y_j \neq 0, \quad 1 \leq i \leq m, \quad 1 \leq j \leq n$$

- Special case of Cauchy matrix: Hilbert matrix
  - $x_i - y_j = i + j - 1$
- Every sub-matrix of a Cauchy matrix is a Cauchy matrix

## Size of the Code

❖ All (N,K) coding can be expressed using encoding and decoding matrices

➤ What's the size of  $c_i$ ?

■ Assume regular  $*$ ,  $+$

➤  $d_i$  is 8 bits


■  $c_i$  will be 18 bits

➤  $d_i$  is 32 bits

■  $c_i$  will be 66 bits

➤  $d_i$  is  $b$  bits

$\Rightarrow c_i$  will be  $b*2 + \log n$  bits



## Size of Code

❖ Code size should be the same as data size

■ Have a longer code than the data size is waste of space

➤ Field

■ Closure for addition and multiplication

• If  $x, y \in F$ , then  $x + y, x * y \in F$

■ Existence of additive & multiplicative inverse in F

• If  $x \in F$ , then  $\exists y \in F$ , s.t.  $x + y = 0$  (or  $x * y = 1$ )

■ Coding computation on a field is reversible

➤ Choice of field for RS computation

■ Integer is a field, but data size has problem

■ Prime field:  $a + b \Rightarrow (a + b) \% p$ ;  $a * b \Rightarrow (a * b) \% p$

• Prime will not be  $2^n \Rightarrow$  One bit longer in code size than in data size

■ Galois' Field: Final choice in coding

## Galois' Field Arithmetic

### ❖ Galois field with $2^N$ elements

- Consider an irreducible polynomial  $P$ 
  - $P = a_0 x^N + a_1 x^{N-1} + \dots + a_{N-1} x + a_N$ 
    - $a_i$ , for all  $i$ , is in  $\{0, 1\}$
  - mod  $P$  is a field
    - (Similar to the concept of mod a prime number)
- Computation
  - If  $p_1$  and  $p_2$  are in  $P$ , then  $p_1 + p_2$  and  $p_1 * p_2$  are in  $P$ 
    - $(p_1 + p_2) \bmod P, (p_1 * p_2) \bmod P$
- Use Galois field in RS computation
  - Map polynomial to the binary numbers
  - $+$  becomes XOR and  $*$  can be done by table lookup

## Galois' Field Arithmetic

### ❖ Galois field with 4 elements

- Consider an irreducible polynomial  $x^2 + x + 1$ 
  - $x^2 + x + 1 = 0 \Rightarrow x^2 = -(x + 1) = x + 1$ 
    - In  $Z_2$ ,  $-1 = -1 + 2 = 1$
  - $(1+x)^2 = 1 + 2x + x^2 = x$
  - $(1+x) * x = x + x^2 = -1 = 1 \Rightarrow x$  and  $(1+x)$  inverse to each other
    - In  $Z_2$ ,  $2x = 0$
  - Map to the 2-bit binary numbers
    - $0 + 0x \Rightarrow 00; 1 + 0x \Rightarrow 10; 0 + 1x \Rightarrow 01; 1 + x \Rightarrow 11$

## Erasure Code in Storage Systems

❖ Erasure code is used quite often in cloud storage

- Microsoft Azure
- GFS II
- Facebook HDFS


➤ Mostly used for redundancy

- Focus on improving recovery speed, i.e., when a share is lost, how to rebuild it efficiently
- If we reconstruct the data, the share can be recomputed, but can we do better?
- All the systems above use an improved method, focusing on recovery speed

## Erasure Code in GFS II


❖ Plain RS 6+3 code in Google GFS II

➤  $(N, K) = (9, 3)$



## Erasure Code in WAS

### ❖ Plain RS 12+4 code



### ❖ Reliability, redundancy, repair cost

#### ➤ 6+3

- Recovery ratio =  $3/9 = 0.33$
- Redundancy =  $3/6 = 0.5$
- Repair cost = 6 shares transferred


#### ➤ 12+4

- Recovery ratio =  $4/16 = 0.25$
- Redundancy =  $4/12 = .33$
- Repair cost = 12 shares transferred  $\Rightarrow$  very expensive!!!

## Erasure Code in WAS


### ❖ Azure's erasure code

#### ➤ Hierarchical code



#### ➤ Handling 1 failure or 2 failures in different groups


- Require 6 share transfers



## Erasure Code in WAS

### ❖ Azure's erasure code


- Handling 2 failures in the same group
  - Require 12 share transfers in either way



## Erasure Code in WAS

### ❖ Azure's erasure code

- Handling 3 failures




## Erasure Code in WAS


### ❖ Azure's erasure code

#### ➤ May handle some of the 4 failure cases

- As long as failures occur in both groups, all four redundant shares are useful




- If failures are all in one group, ( $p_y$  cannot be made use of in the example)  $\Rightarrow$  cannot recover, but chance of this is relatively low



## Windows Azure Storage

### ❖ File structure

- File is considered as a stream, and is “append only”
- Extent: unit of placement/replication
- Block: regular file blocks, each extent has a set of blocks
- Stream means append only  $\Rightarrow$ 
  - Only one extent open for append, the rest are “sealed”
  - Replication for unsealed, erasure coding for “cold” sealed






## Erasure Code in Facebook HDFS-Xorbas

### ❖ Locally repairable code (LRC)

- A type of regenerate code
- A type of hierarchical regenerate code
  - A simplified version
  - Give a new name due to the naming game
- Facebook HDFS Xorbas

## Erasure Code in Facebook HDFS-Xorbas



- $S_1 = C_1X_1 + C_2X_2 + C_3X_3 + C_4X_4 + C_5X_5$
- $S_2 = C_6X_6 + C_7X_7 + C_8X_8 + C_9X_9 + C_{10}X_{10}$

- $S_3 = F_1P_1 + F_2P_2 + F_3P_3 + F_4P_4$   
 $= G_1S_1 + G_2S_2$
- S<sub>3</sub> is implied, not actually stored

## Erasure Code in Facebook HDFS-Xorbas

### ❖ Handling failures

- If at most one failure in each group  $\Rightarrow$  local repair
  - If  $P_1$  fails, compute  $S_3 = G_1S_1 + G_2S_2$ , and then derive  $P_1$  from  $S_3$
- If more than one failure in a group  $\Rightarrow$  require RS repair
- No problem in handling any 6 failures
  - No problem with 4 failures
  - If all 5 blocks in a group fail, recover from 1 local + 4 RS
  - The 6th failure has to be in a different group
- If only 3 RS blocks are used
  - Same as the case in Azure erasure code, may not be able to recover 5 failures with 5 redundant blocks

## References

### ❖ References

- RAID
  - A case for redundant arrays of inexpensive disks (RAID)
  - M. Blaum, J. Brady, J. Bruck and J. Menon, "EVENODD: An efficient scheme for tolerating double disk failures in RAID architectures," IEEE Transactions on Computing, Vol. 44, No. 2, Feb. 1995, pp. 192-202.
  - Jay White, Chris Lueth, Jonathan Bell, "RAID-DP: NetApp Implementation of Double-Parity RAID for Data Protection" NetApp.com, March 2003.
  - C. Huang and L. Xu, "STAR: An efficient coding scheme for correcting triple storage node failures," Usenix Conference on File and Storage Technologies (FAST), December, 2005, pp. 197-210.

## References

### ❖ References

- Reed Solomon
  - I. S. Reed and G. Solomon, “Polynomial codes over certain finite fields,” Journal of the Society for Industrial and Applied Mathematics, 8, 1960, pp. 300-304.
  - Vandermonde matrix: [http://en.wikipedia.org/wiki/Vandermonde\\_matrix](http://en.wikipedia.org/wiki/Vandermonde_matrix)
  - Cauchy Reed-Solomon: <http://planetmath.org/CauchyMatrix.html>
- Azure erasure storage
  - C. Huang, H. Simitci, Y. Xu, A. Ogus, B. Calder, P. Gopalan, J. Li, S. Yekhanin, “Erasure coding in Windows Azure storage,” USENIX ATC, June 2012
- Facebook HDFS-Xorbas
  - M. Sathiamoorthy, M. Asteris, D. Papailiopoulos, A.G. Dimakis, R. Vadali, S. Chen, D. Borthakur. “XORing elephants: Novel erasure codes for big data,” VLDB 2013, pp. 325-336