

Big Data Computing Models

Map-Reduce

What is Map-Reduce?

❖ Parallel Programming Models

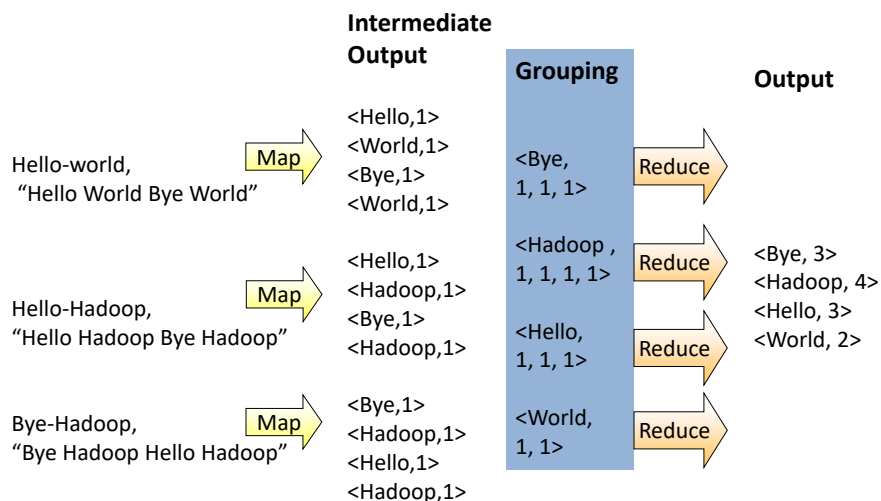
- Message passing
 - Independent tasks encapsulating local data, interact by exchanging messages
- Shared memory
 - Tasks share a common address space and interact by reading and writing this space asynchronously
- Data parallelization
 - Perform the same operations to a set of data
 - Also referred to as “Embarrassingly parallel”

❖ Map reduce is a form of data parallelization

What is Map-Reduce?

- ❖ Map-reduce makes parallel programming easier
 - Coarse grained data parallel computing
 - User only writes a mapper and a reducer
 - MapReduce system takes care of the rest
- ❖ Use functional programming model
 - Function as an argument
 - Map/reduce (function (list))
 - Map (square (1 2 3 4 5)) = (1 4 9 16 25)
 - Reduce (+ (1 2 3 4 5)) = 15

Example -- Word Count



<https://hadoop.apache.org/docs/current/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>

Example -- Word Count

❖ Map and reduce program

```
map (key, value):  
    // key: document name; value: text of document  
    for each word w in value: emit(w, 1)  
reduce (key, values):  
    // key: a word; values: the count of the word  
    result = 0  
    for each v in values: result += v  
    emit (key, result)
```

*emit in java mapreduce:
setup(Context context);
context.write (...);*

➤ In a large document

■ Input to word count

- Key: line number; value: text in the line

https://hadoop.apache.org/docs/current/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html#Example:_WordCount_v2.0

Map-Reduce Process

❖ Input

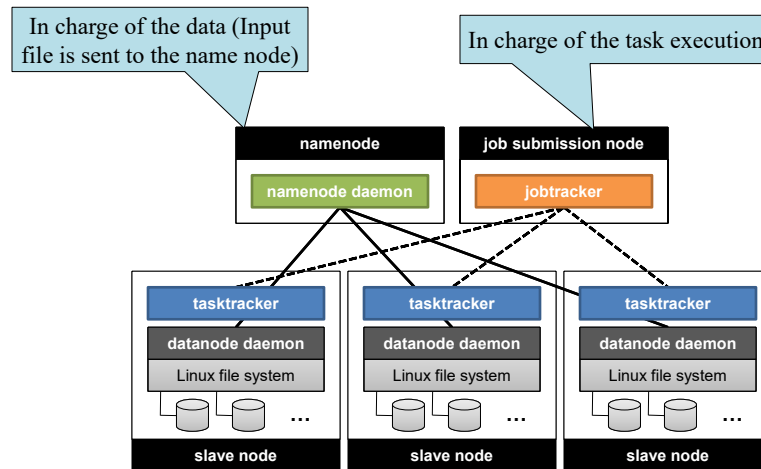
- A list of key/value pairs: list(key, val)

❖ User supplies two functions:

- Map-function (list(key, val)) → list(key', val')
 - list(key', val') is a list of intermediate key-value pairs generated from mapping function
 - Original list and new list may not be in the same size
 - One key' may have many different values
- Reduce-function (list(key_i', val')) → v_i, for all i

❖ The execution framework handles everything else

Hadoop Map-Reduce Framework



Hadoop Map-Reduce Framework

❖ Master-Worker structure

- MapReduce job is sent to the master node JobTracker
- JobTracker
 - Coordinate the execution of the tasks
 - Distributes the mapper task to TaskTrackers
 - Starts the TaskTrackers with the tasks
 - Wait for TaskTrackers' results
 - » If some TaskTrakers have failed, redistribute the missed tasks
 - Determines reduce task allocation and sends TaskTrackers the tasks
 - » Give reducer the intermediate file locations and sizes to reducers
 - Wait for TaskTrackers' results
 - » If some TaskTrakers have failed, redistribute the missed tasks
 - Also periodically ping the TaskTrakers to keep track of their status

Hadoop Map-Reduce Framework

❖ Master-Worker structure


➤ TaskTracker -- Map phase

- Creates map task instances and runs them
 - Intermediate results are stored in a circular buffer in memory
 - When buffer is full, spill to local disk
- Combiner combines intermediate results, sort them by key
 - Combiner mostly uses a hash function
- Sends locations and sizes of intermediate files to the JobTracker

➤ TaskTracker -- Reduce phase

- Merger obtains intermediate results from mappers and merge them
- Creates reduce task instances and runs them
- Send the results to the JobTracker

Hadoop Map-Reduce Framework



Hadoop Map-Reduce Framework

❖ Master-Worker structure

- Final output
 - TaskTracker responds to job issuer about the completion of the task
- Final output may go for next map-reduce task
 - In this case, programmer does not combine the output files into one
 - Simply passes the output files (same name) as input to another MapReduce call
 - Cannot have shuffling from reduce to next round map
- Map-reduce input and output
 - Always files (or a nosql-DB) that are partitioned into blocks and distributed over multiple nodes

Hadoop Map-Reduce Framework

❖ Synchronization

- Reduce phase does not start till map phase is fully done
 - Choose to avoid complex issues in starting some reducers early
 - Use the idle workers to help speed up the remaining map phase

❖ When most subtasks (map or reduce) are done

- Slow workers significantly delay completion time
 - Bad disks w/ soft errors transfer data slowly
 - Worker failure
- Use multiple workers to execute the remaining tasks to ensure completion
 - Anyway these workers are idle

Hadoop Map-Reduce Framework

❖ Fault tolerance

- Always assign one subtask to three workers
 - Even if all replicas fail \Rightarrow No results are generated \Rightarrow freed-up workers will re-execute the task
 - During a massive failure, lost 1600/1800 nodes, task finished fine
- Do not handle master failure at this point

❖ When tasks fail

- Map/Reduce functions sometimes fail for particular inputs
- The worker sends the problem to the master node
 - Include sequence number of record being processed
- If master sees two failures for the same record
 - Give up the subtask and stop all worker replicas for the same record

Hadoop Map-Reduce Framework

❖ Mini Reducer

- Programmer may specify a “Combiner” function that does data merging on a single node
- Typically, the same code is used to implement both the combiner and the reduce function
- Example
 - Word count: Compute the cumulative count at each node
 - In order for this to work, the reduce function should be commutative and associative

Hadoop Map-Reduce Framework

❖ Important performance booster

- Compress the mapper output to reduce the communication cost during shuffling
 - Tradeoff between compression/decompression cost and communication cost
 - Hadoop offers compression option

Evolution of Map Reduce Frameworks

What's New in MapReduce?

❖ Data parallel computation

- In 80s, MIT developed the Connection Machine
 - Also called the thinking machine
 - The corresponding parallel programming paradigm is called data parallel programming (or more generally, the **SIMD** programming model)
 - Map (instructions (list))
- In 80s, Teradata and Gamma project developed a new parallel database paradigm
 - Partition database row-wise and distribute partitions to nodes
 - Perform database operations in parallel in a cluster
- What is new in MapReduce?

SIMD Programming

❖ Programming model

How is MapReduce
different from SIMD

- Define the basic data types
 - Can be scalar data type, or data structures
- Define a special **array** of a basic data type
 - The special array is distributed over the computing nodes
 - E.g., in connection machine, it is called “shape”
 - E.g., shape data[N], sqdata[N];
- Computation is **instruction based**
 - Computation in MapReduce is task based, more coarse-grained
 - E.g., Map: sqdata[i] = square (data[i]);
 - E.g., Reduce: sum = +sqdata[i];
- Communication is specified based on array indices
 - Communication in MapReduce is specified by keys
 - E.g., data[i] = data[j] + data[j+1]

Issues in Hadoop -- No Loop

❖ MapReduce does not support iterative computing

- And, fixed M-R-M-R flow (how about M-R-R-M)

❖ Multiple stages of MapReduce

- User has to match I/O data files
 - Create two JobConf objects, first job has the normal input, but output to temp, for example, and second job has temp as input and produces the normal output file
- Or define job dependencies through jobControl object
 - `Job job1 = new Job(...); Job job2 = new Job(...);`
 - `JobControl jbcntrl = new JobControl("jbcntrl");`
`jbcntrl.addJob(job1); jbcntrl.addJob(job2);`
`job2.addDependingJob(job1);`
 - `jbcntrl.run();`


Loop Example -- Page Ranking

❖ A common task for Google

❖ Definition

- Given a page x , with inlink from pages t_1, t_2, \dots, t_n
 - $C(t_i)$: out degree of page t_i
 - α : probability of random jump
 - N : the total number of nodes in the graph



$$PR(x) = \alpha \left(\frac{1}{N} \right) + (1 - \alpha) \sum_{i=1}^n \frac{PR(t_i)}{C(t_i)}$$





Example -- Page Ranking

$$PR(x) = \alpha \left(\frac{1}{N} \right) + (1 - \alpha) \sum_{i=1}^n \frac{PR(t_i)}{C(t_i)}$$

Iteration i



Iteration i+1



Page Ranking in Hadoop

❖ Page rank computation

➤ Stage 0: Process each page x

■ Map phase:

- Initiate x 's page rank, $x.rank$, to 1
- Process page content, compute $x.outlist$ (the list of x 's outgoing links), and compute $x.\#links$ (the total number of outgoing links in x)
- Compute $credit = x.rank / x.\#links$
- For each url in $x.outlist$, create an item $\langle key=url, val=credit \rangle$

■ Reduce phase:

- Each reducer will receive $\langle key=y, value=credit \rangle$ from all its inlinks
- Compute the sum of credits and assign it to $y.rank$

➤ Note: After reduce, there is no shuffle and exchange

- $y.rank$ is the output and is to be used in the next iteration

Page Ranking in Hadoop

❖ Page rank computation

➤ Stage i :

- Map phase: (only distribute each rank received)
 - Need $x.rank$, $x.outlist$ and $x.\#links$
 - » $x.rank$ is from the previous round
 - » $x.outlist$ and $x.\#links$ has been generated in the first round
 - Compute $credit = x.rank / x.\#links$
 - For each url in $x.outlist$, create an item $\langle key=url, val=credit \rangle$
- Reduce phase
 - Same as Stage 0

➤ Termination check

- No change of page rank for all pages

Issues in Hadoop -- No Loop

❖ MapReduce does not support iterative computing

- User has to do it outside MapReduce \Rightarrow Inefficient solution
Why inefficient?

❖ Haloop MapReduce

➤ Provide new constructs for expressing loops

- $D_{i+1} = f(D_i, L)$
- D_{i+1} : reducer output of the i -th iteration, which will be the mapper input of the $(i+1)$ -th iteration
 - In Hadoop, they would be stored in HDFS Inefficient!
- L : invariant data, does not change over the iterations
 - In Hadoop, they were in memory and would be gone after one MR Inefficient!

➤ Termination

- Programmer defines the termination condition
 - E.g., $D_{i+1} = D_i$, or $|D_{i+1} - D_i| < \epsilon$, or reach a certain bound, or ...

Issues in Hadoop -- No Loop

❖ Haloop MapReduce

- Workers maintain cache and local files
 - Cache the reducer outputs (D_{i-1}), which will be the input of the mapper in the next iteration *What are these in the page rank program?*
 - Cache the invariant data (L) in the first iteration to allow reusing in later iterations *What are these in the page rank program?*
 - Perform local termination evaluation
 - Master sends the evaluation criteria with the reduce tasks
 - Send the evaluation result to Master
- Master
 - Repeats the map-reduce program
 - Coordinate the check of termination condition
 - Master acts as the single node reducer

Issues in Hadoop -- Data Placement

❖ Data placement problem

- Sample problem: Join
 - Get the list of names and emails of customers with transaction \$ amount > \$1000 (for targeted customer retention)
 - Customer info table
 - » Sorted by C.custId
 - Sales transaction table
 - » Sorted by transaction time
- HDFS
 - Does not allow users to specify how the input files should be partitioned and placed
- Improved systems: Hadoop++, Liah, CoHadoop, ...


```
SELECT C.name, C.email
FROM Customers C, Sales S
WHERE C.custId = S.custId
AND S.amount > 1000
AND S.date BETWEEN
"12/1/15" AND "12/25/15";
```

Issues in Hadoop -- Data Placement

❖ Hadoop

- Data are partitioned according to the original input file order and placed on multiple nodes at random
- Need to first perform MapReduce to align the keys in the two files

```
SELECT C.name, C.email
FROM Customers C, Sales S
WHERE C.custId = S.custId
AND S.amount > 1000
AND S.date BETWEEN
"12/1/15" AND "12/25/15";
```




Issues in Hadoop -- Data Placement

❖ Hadoop++

- Data are pre-partitioned into splits
 - User can determine which key to use for pre-partitioning (e.g. custId in the example)
 - A Trojan index describes each split
 - Tree based indexing of splits (key range, #records, ...)
 - Stored on disk, allow easy cache in memory
 - Not available in Hadoop

```
SELECT C.name, C.email
FROM Customers C, Sales S
WHERE C.custId = S.custId
AND S.amount > 1000
AND S.date BETWEEN
"12/1/15" AND "12/25/15";
```



After prepartition:
Data with the same keys
are placed in the same split
(e.g., custID given by user)

Issues in Hadoop -- Data Placement

❖ Hadoop++


➤ For PageRank example

- Many links from a page are in the same domain
⇒ Partition the links by domain name instead of by the entire url
⇒ Can greatly reduce shuffling cost for each iteration

Issues in Hadoop -- Data Placement

❖ CoHadoop

- Data are pre-partitioned based on the key (like Hadoop++)
- Also support co-location for multiple datasets
 - Use a locator group concept
 - A data object can have a locator number
 - Multiple data objects with the same locator number are in one locator group and are co-located (blocks with the same key range are placed on the same node)



Issues in Hadoop

❖ Minor issue

- Does not perform well on heterogeneous clusters
 - Hadoop attempts to distribute loads evenly
 - E.g., divide a data file evenly across the nodes
 - If the physical nodes have different computing capacity
 - The scheduler will reschedule the tasks for slow nodes to fast nodes
 - Require data movement from slow nodes to fast nodes
 - If partitioning of the data can go by the computing capacity
 - All nodes will progress in a similar rate
- Easy to fix
 - E.g., using virtual nodes to make the system homogeneous
- This can also be considered as a data placement problem

Issues in Hadoop

❖ Require detailed programming

- Does not support SQL ⇒ Need detailed coding for some standard DB operations
 - E.g., select, join, group-by, ...
 - Complaints mainly from the DB community
- Solution
 - Provide standard DB operations on top of MapReduce
 - E.g., Hive, Pig, Shark, ...
 - Hive (HQL), originally by Facebook, then shifted to Apache
 - Pig (Pig Latin), originally by Yahoo, then shifted to Apache
 - Shark: Hive on Spark

Pig

❖ Pig: <http://incubator.apache.org/pig>

- Developed by Yahoo!, now open source
- Use Pig Latin, a dataflow language
 - Why not SQL?
 - Too complex to achieve efficient large data processing
 - Eliminate the non-common constructs, suitable for almost all commonly used queries
- Support a fully-nestable data model
 - More natural and flexible than relational DB's flat tuples
 - Avoids expensive joins as much as possible

■ E.g.,

$\left(\text{yahoo}, \left\{ \begin{array}{l} \text{finance} \\ \text{email} \\ \text{news} \end{array} \right\} \right)$

Similar to the data model of some NoSQL DB systems

Pig -- Example

❖ Example

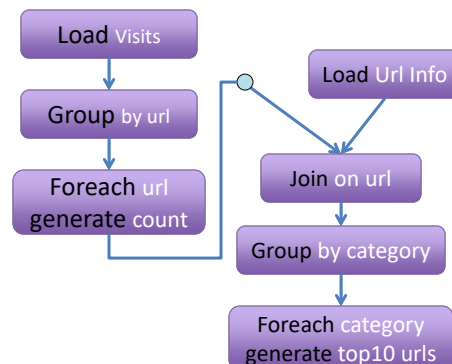
- Given two tables: Visits, url-Info
 - Originally: Visits: ordered by time; url-info: ordered by url
- Find the top 10 most visited pages in each category

Visits

User	Url	Time
Amy	cnn.com	8:00
Amy	bbc.com	10:00
Amy	flickr.com	10:05
Fred	cnn.com	12:00

url-Info

Url	Category	PageRank
cnn.com	News	0.9
bbc.com	News	0.8
flickr.com	Photos	0.7
espn.com	Sports	0.9



Pig -- Example

❖ Example



```

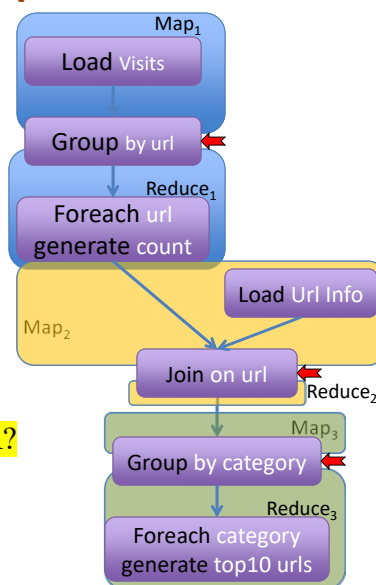
visits = load '/data/visits' as (user, url, time);
gVisits = group visits by url;
visitCounts = foreach gVisits
    generate url, count(visits);
urlInfo = load '/data/urlInfo' as
    (url, category, pRank);
visitCounts = join visitCounts by url,
    urlInfo by url;
gCategories = group visitCounts by category;
topUrls = foreach gCategories
    generate top(visits,10);
store topUrls into '/data/topUrls';
  
```

* top, count are user defined functions

Pig -- Example

❖ Compile pig into map reduce

- Each group and join forms a map-reduce boundary
 - Has to be done by both map and reduce phases
- Other operations pipelined into map-reduce phases
 - E.g., foreach can be done in map or in reduce
- What is the key for each M to R?



Hadoop Variants So Far

❖ Hive, Pig

- Provide a high level language ⇒ Easier to program
- Provide tools to transform the high level code to map reduce code
- These systems are built on top of Hadoop ⇒ Still have the same problems as in Hadoop

❖ Haloop, Hadoop++

- Manipulate data placement
- Still on top of Hadoop

❖ How about completely change Hadoop

- Spark

Spark

❖ Developed

- By Matei Zaharia at Berkely as his PhD work
- Then open sourced by Apache

❖ Major concept

- RDD: **Resilient Distributed Dataset**
 - In Hadoop, datasets are stored as files on disk (HDFS)
 - Input/output datasets for a job have to be disk files
 - Data used across jobs are subsequently disk files
 - In Spark, each dataset is represented as an RDD
 - A Spark operation takes input RDDs and create output RDDs
 - RDD is in memory, distributed over multiple nodes
 - Persistent RDDs also have corresponding HDFS copies on disk
 - » Asynchronously written to disk

Spark RDD

❖ Persist RDD

- Call “persist” or “cache” methods to mark the RDD
 - By default, the RDD will be removed after the workflow
 - Can also call “unpersist” method to remove it
- The persist RDD will be retained in memory
 - ⇒ Allows future operations to reuse the RDD
 - ⇒ Much faster than disk accesses
 - ⇒ Even if no loop support ⇒ No performance problem
 - If memory is abundant, 30-100 folds of speedup
- Can choose different storage levels
 - (MEMORY_ONLY, DISK_ONLY, MEMORY_AND_DISK, etc.)
 - MEMORY_AND_DISK
 - Move some RDD partitions to disk when memory is full
 - Use LRU policy

Spark RDD

❖ RDD serialization

- Hadoop persistent objects are writable
 - Has de/serialization costs
- RDD avoids serialization (in memory)
 - Can be serialized as desired
 - Spark supports both java and kryo serialization (kryo does not support all data types, but is much faster)

❖ No replication

- Spark gives up RDD replication
 - RDD is immutable once created, read only
 - RDD uses logging to achieve resilience

Spark RDD

❖ RDD resilience

- No replication in RDD
 - Replication is not scalable in both computation performance and storage requirement (persistent RDDs are stored in HDFS)
- Lineage: log the operations performed on RDDs
 - Represented by DAG
 - Nodes in DAG: RDDs; Edges in DAG: Spark operations
- Data recovery: RDD + lineage
 - A lost partition of an RDD can be recomputed automatically
 - Original RDD would have been an HDFS file
 - Operations has to be coarse grained to avoid logging overhead
- Checkpoint RDDs
 - Prevent from having long lineage chains

Spark Operations

❖ Map-Reduce operations

- map (func)
- reduce(func)
- collect()
 - Similar to reduce, with one reducer
 - Returns all the elements of the dataset at the driver (master)
- count()
 - Not available in Hadoop
 - Returns the number of elements in the dataset

Spark Operations

❖ Add some DB operations

- join(another-dataset)
- groupby(column)
- filter(func)
 - Similar to select (new dataset is a subset of source dataset)
 - func defines the selection criteria (func returns true/false)

❖ Add some Set operations

- union(another-dataset)
- intersection(another-dataset)
- distinct()
 - Returns a new dataset that contains the distinct elements of the source dataset (remove duplicates)

Spark Operations

❖ Task

- Unit operation

❖ Stage:

- A sequence of tasks that does not have shuffle and exchange in between (e.g., map, filter, union, intersection)
- Does not enforce synchronous operations
 - Workers can run asynchronously for operations in a stage
- Otherwise, same as Hadoop, synchronous execution

Spark Operations

❖ In Hadoop

$N = \text{\#HDFS blocks for the dataset}$; $M = \text{\#keys from mapper}$

- N mappers, M reducers, C cores
- Output of each mapper is sorted and divided into M objects
- There are $N*M$ objects to be shuffled

❖ In Spark

- Initially, the same as Hadoop
- Later: Output of mappers on the same core are merged
 - There are only $C*M$ objects to be shuffled
 - This optimization is hardware dependent, but does not burden the user (since shuffle phase is anyway hidden from the user)
 - Note: Mini combiner in Hadoop is different, only merge data entries generated from the same mapper

Spark Data Placement

❖ RDD partitioning

- Default: by hashing (HashPartitioner)
- User can define data partitioning and co-location

❖ Partitioner object

- Can specify the partition rules
 - Similar to Hadoop++
- Each RDD can specify an optional Partitioner object
 - Two RDDs share the same Partitioner are co-located
 - Similar to locator in CoHadoop
 - Partitioner does not have to have partition rules, just for co-location

Spark

❖ Cluster architecture is the same as Hadoop

- Driver = Master in Hadoop
 - Defines and invokes actions on RDDs
 - Tracks the RDDs' lineage
- Workers
 - Store RDD partitions
 - Perform RDD transformations

Spark



❖ Difference from Hadoop

- Storage: Spark has in-memory store
 - Spark uses in-memory persistent storage
 - Also have copies on HDFS for persistent RDDs
- No RDD replication
 - Resilience is via lineage logging
- Computation: Spark has a more flexible model
 - Support loop, and can have any M-R combinations
 - Provide more operators, rather than just map and reduce
- Support co-location and placement specification
- Overall
 - ⇒ Fix many problems in Hadoop
 - ⇒ Offer the RDD solution ⇒ Unique in MapReduce evolution

Spark Performance

❖ Page rank example

- Links: RDD indexed by url
 - Value: outgoing links
- Ranks: RDD indexed by url
 - Value: page ranking
- Computation requires repeated join of Links+Ranks



References

❖ MapReduce

- MapReduce: Simplified data processing on large clusters
- Parallel data processing with MapReduce: A survey

❖ Improvements

- HaLoop: Efficient iterative data processing on large clusters
- CoHadoop: Flexible Data Placement and Its Exploitation in Hadoop
- Building a highlevel dataflow system on top of MapReduce: The Pig experience
- Hive: A warehousing solution over a MapReduce framework

❖ Spark

- Spark: Cluster computing with working sets
- Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing